

ModeSwitch

The right content. For your now.

"YouTube makes me decide what I want when I opened it because I didn't want to think in the first place."

"A person has so many personalities in a day — why have an algorithm that's updated based on what you saw last night?"



THE PROBLEM

YouTube has a rich picture of what you like. It doesn't explicitly know what you're here for.

What YouTube optimises for

Past watch history

Long-term content preferences

Historical engagement patterns

What performed well with similar users

What's missing

What you need from this session

The context you're watching in

Whether you want to think or switch off

What job you're trying to get done tonight

The algorithm is a historian. ModeSwitch adds one missing signal: what you're trying to get from this session.

RESEARCH

Across user conversations, one pattern kept repeating.

Users often open YouTube for a job beyond “watch a video” — they want to decompress, fill time, shift their mood, or have something on in the background.

Decompress without effort

Brain is fried. Don't want decisions.

Fill a small pocket of time

10–15 minutes. Instant engagement.

Change how I feel

Stuck or low. Want a mood shift.

Not be alone while doing something else

Background company while cooking, cleaning, working.

Escape the blankness

Bored but nothing sounds interesting. Don't even know what I want.

These are jobs, not categories. That distinction shapes the solution.

Evidence: secondary research + qualitative user conversations. Estimates labeled where applicable.

FROM RESEARCH TO PRODUCT

How the research created the product decisions.

RESEARCH PATTERN	PRODUCT IMPLICATION
<i>"I don't want to think when I open YouTube"</i>	Don't make users configure intent upfront. Value before configuration.
<i>"I know roughly what I want, but not what to watch"</i>	Capture intent before content selection, not after.
<i>"Sometimes I just want something on in the background"</i>	Treat background viewing as a legitimate session intent, not a fallback.
<i>"This phase of watching doesn't reflect my normal interests"</i>	Isolate temporary modes — don't let missions reshape permanent recommendations.
<i>"I don't always know what I want until I see options"</i>	Suggestions must be optional. "Not now" is a first-class behavior.
<i>"My mood is hard to describe in words"</i>	Infer context from behavior — don't ask users to label their emotional state.

This table is the product brief. Every feature decision in ModeSwitch traces back to one of these rows.

THE OPPORTUNITY

There is substantial selection friction. Research suggests context mismatch is a contributor.

122M

Daily active users

Global Media Insight, 2024

70%

Of watch time driven by the algorithm

YouTube CPO via MIT Tech Review

3–8%

Homepage click-through rate

Miraflow, 2026 · Secondary research

55%

Of video clicks drop off within 60 seconds

Retention Rabbit, 2025 · Secondary research

Intent ambiguity

Homepage loads, no intent signal, user scrolls

1–2+ min



Selection failure

Click happens, mood mismatch reveals itself
Drop-off within 60 sec



Successful start

Content matches context, user settles in
Session begins naturally



Satisfied session

Content continues to fit intent, session proceeds
No implied time target

↑ ModeSwitch intervenes here — before the click, not after.

Goal is not to maximize session length. Goal is to reduce the time between opening YouTube and feeling satisfied with what's playing.

Evidence: secondary research + qualitative user conversations. Estimates labeled where applicable.

THE FEATURE

ModeSwitch is not a recommendation fix. It's an intent layer.

"ModeSwitch helps YouTube understand your intent in the moment — not just your interests over time."

What it's not

A profile switcher

A playlist

A mood filter

A recommendation redesign

What it is

A session entry layer

An explicit intent signal given to the algorithm

A context-aware homepage reshaper

A system that learns from how you actually watch

Less searching. Faster state change.

THE CONTEXT · WHY TV

TV is where session intent matters most.

TV represents a more committed viewing context than micro-moment mobile browsing. The intent to settle in is already present. ModeSwitch captures it.

Higher session commitment

When a user opens YouTube on TV, they are typically ready to watch something, not browse between tasks.

Lean-back environment

TV viewing is less interruptible. Getting into the right content faster has more impact when the alternative is extended remote-scrolling.

Remote-first interaction

TV navigation requires fewer, larger choices. Design constraint: value must be apparent within one glance and ≤ 3 remote interactions.

These principles apply equally to mobile. TV is the starting point because the context-intent mismatch is most visible there.

THE FEATURE · PRINCIPLES

Four principles. Every design decision is held to these.

PRINCIPLE 01

Intent, not identity

Understand what the user needs now — not who they are overall. A mode is a session signal, not a personality type.

PRINCIPLE 02

Suggest, don't force

Inference should reduce effort, not create another decision. "Not now" is always valid and the system learns from it.

PRINCIPLE 03

Lens, not silo

Modes influence ranking without fragmenting the user's profile. High-confidence preferences can flow to the global taste graph.

PRINCIPLE 04

Infer context, not emotion

Never claim to know how someone feels. "Want something easy tonight?" is right. "Looks like you're stressed" is not. Infer viewing intent, never emotional state.

THE FEATURE · PRODUCT VISION

The product vision: one system, three layers.

CORE

Suggested Mode

YouTube reads time, device, and recent session behavior. Makes a call before you have to. "Want something easy tonight?" One tap. No configuration. This is the layer worth testing first — and the only layer needed to validate the hypothesis.

PERSONALIZATION

Your Modes

Create named modes. Seed with 2–3 channels or topics. YouTube learns from what you actually watch inside each one. Permanent modes are porous — high-confidence preferences can influence your global profile because lasting recommendations should reflect them.

ADVANCED


Quick Modes

Japan trip. Interview prep. New puppy. Fully isolated from global profile and other modes — because the intent is explicitly temporary. The system treats completion as probabilistic, not certain: declining engagement triggers a quiet prompt. Revivable within 30 days.

Starter Modes — a few lightweight defaults on first launch — exist only in onboarding. They are not a permanent layer.

THE FEATURE · ONBOARDING

Value before configuration.




The user experiences ModeSwitch before being asked to name, build, or configure anything. Configuration is optional and comes after the value is felt.

This resolves the core tension: the feature promises to reduce thinking, so its onboarding must not create new thinking.

THE FEATURE · SYSTEM

ModeSwitch is a state machine, not a UI feature.




■ User state ■ Product interface ■ System intelligence

Optimise for intent clarity at session start. Separate what the user needs now from who they are overall.

THE FEATURE · IN ACTION

What it looks like on TV.



Vista TV is a fictional platform used for demonstration. This concept is designed for YouTube.

Video not playing? Watch it here: youtu.be/ChU_RlpcaR4

METRICS

Adoption tells us if users are using it. Value tells us if it's working. Guardrails tell us if we're breaking anything.

NORTH STAR

% of eligible sessions with an active intent context within the first 10 seconds

Sub-signal · User-selected intent
Sub-signal · Suggested intent accepted
Sub-signal · Continued previous intent

DIRECTIONAL TARGETS

Month 1 · ~10% Month 3 · ~25%
Month 6 · ~40% Month 12 · ~55%

Directional targets, not benchmarks — to be calibrated through experiment.

LEADING INDICATOR

Suggested Mode acceptance · $\geq 30\%$

Are context signals accurate enough that users trust the suggestion?

LEADING INDICATOR

Time-to-first-video, mode-active · meaningful reduction vs. baseline

Is the promise being kept?

LEADING INDICATOR

Mode adoption at 30 days · $\geq 35\%$ still using

Is it becoming a habit?

GUARDRAILS — IF ANY TRIP, ROLLOUT PAUSES

Pre-play abandonment — must not increase platform-wide
New user Day 7 retention — any negative delta triggers review

Overall session watch time — no regression across all users
Recommendation satisfaction — mode-filtered feeds must not reduce quality vs. baseline

TRADEOFFS

The honest version of this doesn't defend the feature. It confronts the real risks.

"You're adding a new decision point at the exact moment we've spent 15 years trying to remove decisions."

Decision point vs. decision removal

Suggested Mode makes the call before the user has to. But ModeSwitch only succeeds if choosing a mode is consistently faster than discovering content through browsing. If it isn't, the feature has failed even if adoption is high.

Overfitting to a rare pain

55% drop-off in 60 seconds and 3-8% homepage CTR suggests substantial friction — but we can't definitively attribute this to context mismatch. Our research points to it as a contributor. That's not the same as proof.

Premature self-awareness

Some sessions, users genuinely don't know what they want. "Not now" is a first-class behavior. The system learns from repeated dismissals in specific contexts and suppresses suggestions accordingly.

Explicit intent vs. algorithmic discovery

A mode gives users control, but excessive adherence can prevent valuable recommendations outside the selected lens. Intent is a ranking signal — not an absolute filter. Modes should bias the feed, not wall it off.

Helpful inference vs. creepy inference

The system infers viewing context, not emotional state. "Want something easy tonight?" is right. "Looks like you're stressed" is not. The design principle: infer intent, never emotion.

WHEN MODESWITCH STAYS OUT OF THE WAY

Resuming an unfinished video — intent already established.

Active search — the user has defined their own intent.

Dismissed 3+ times in the same context — learn and suppress.

ROADMAP · WHAT I'D BUILD

What I would actually build first.

PHASE 1

Test the hypothesis

Suggested Mode only. One context-aware prompt at session start. No mode bar, no creation flow, no Quick Modes. If this alone changes time-to-first-video meaningfully, the hypothesis holds. If it doesn't, the architecture needs to change before a full build begins.



PHASE 2

Build the habit

Persistent user-created modes. Mode bar on homepage. Porous learning model. Users build their own intent layer on top of YouTube.



PHASE 3

Handle missions

Quick Modes with probabilistic lifecycle detection. Temporary intent that doesn't reshape permanent recommendations.

This phasing keeps the product aligned with its own thesis: reduce the amount of thinking required before value is felt.

REFLECTION

This wasn't a writing exercise.

What surprised me

I started believing recommendation quality was YouTube's core problem. Research shifted my thinking toward intent ambiguity. The algorithm may be surfacing good content — but not necessarily the right content for the moment. That reframe changed everything.

What changed during the process

The original concept used fixed preset modes. User conversations pushed toward user-defined modes with system suggestions. That shift moved personalization from the interface to the core of the product — and made ModeSwitch harder to build and more valuable to use.

What I'd test first

Suggested Mode only. One prompt, no configuration, no mode bar. If that single signal changes time-to-first-video meaningfully, the hypothesis holds. If it doesn't, the problem is more complex than intent signaling can solve.

"The success of ModeSwitch is not whether the feature ships. It's whether the hypothesis survives contact with real user behavior."

Sanyukta Jha · [linkedin.com/in/sanyuktajha](https://www.linkedin.com/in/sanyuktajha)